

[CSED211] Introduction to Computer Software Systems

Lab 6: Cache Lab

Dowon Son



CAOS
COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

2024.11.14

Today's Agenda



- Background
- Cache Lab Part A: Building Cache Simulator
 - Valgrind
 - File I/O APIs
 - Dynamic Allocation & Deallocation
 - Parsing Command Line Options
- Cache Lab Part B: Efficient Matrix Transpose
 - Hit Ratio
 - Matrix Multiplication

Today's Agenda

- Background
- Cache Lab Part A: Building Cache Simulator
 - Valgrind
 - File I/O APIs
 - Dynamic Allocation & Deallocation
 - Parsing Command Line Options
- Cache Lab Part B: Efficient Matrix Transpose
 - Hit Ratio
 - Matrix Multiplication

Background: Locality

- **Principle of Locality:** Programs tend to use data and instructions with addresses **near or equal** to those they have **used recently**
- **Spatial locality:**
 - Items with nearby addresses tend to be referenced close together in time
- **Temporal locality:**
 - Recently referenced items are likely to be referenced again in the near future



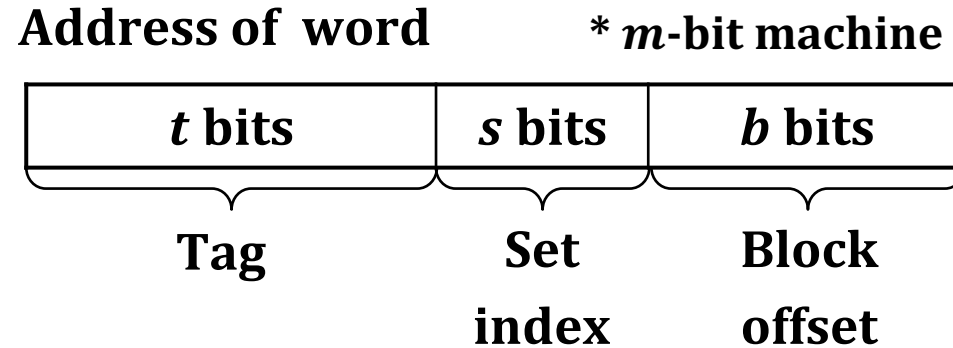
Background: Cache Memory

- **Cache memory** is small, fast SRAM-based memory
 - **Automatically managed** in hardware
 - Holds frequently accessed blocks of main memory
- CPU first looks for data in caches (e.g., L1, L2, and L3), then in main memory

Background: Cache Memory (Cont.)

- Cache set
 - Cache memory consists of $S (= 2^s)$ cache sets
- Cache line
 - Each cache set consists of E (associativity) cache lines
 - Each cache line consists of valid bit, tag, and (2^b-byte) data block
- Valid bit: Shows the cache line is **valid** or **not**
- Tag: Compares the tag of the cache line with the currently accessed memory address to check **for a match**
- Data block: Saves the **loaded data**

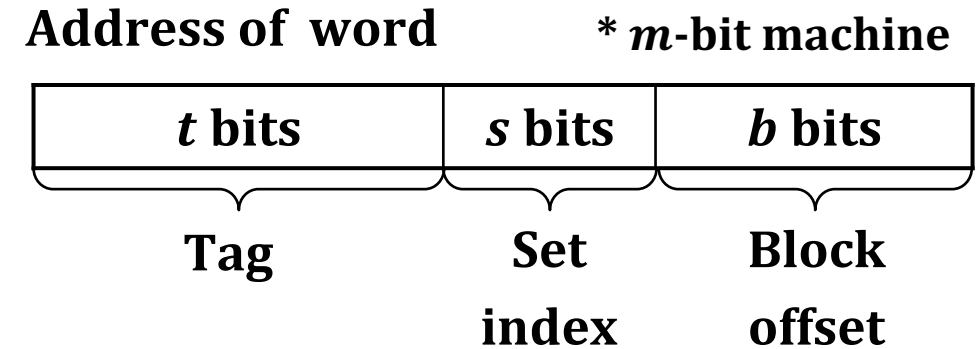
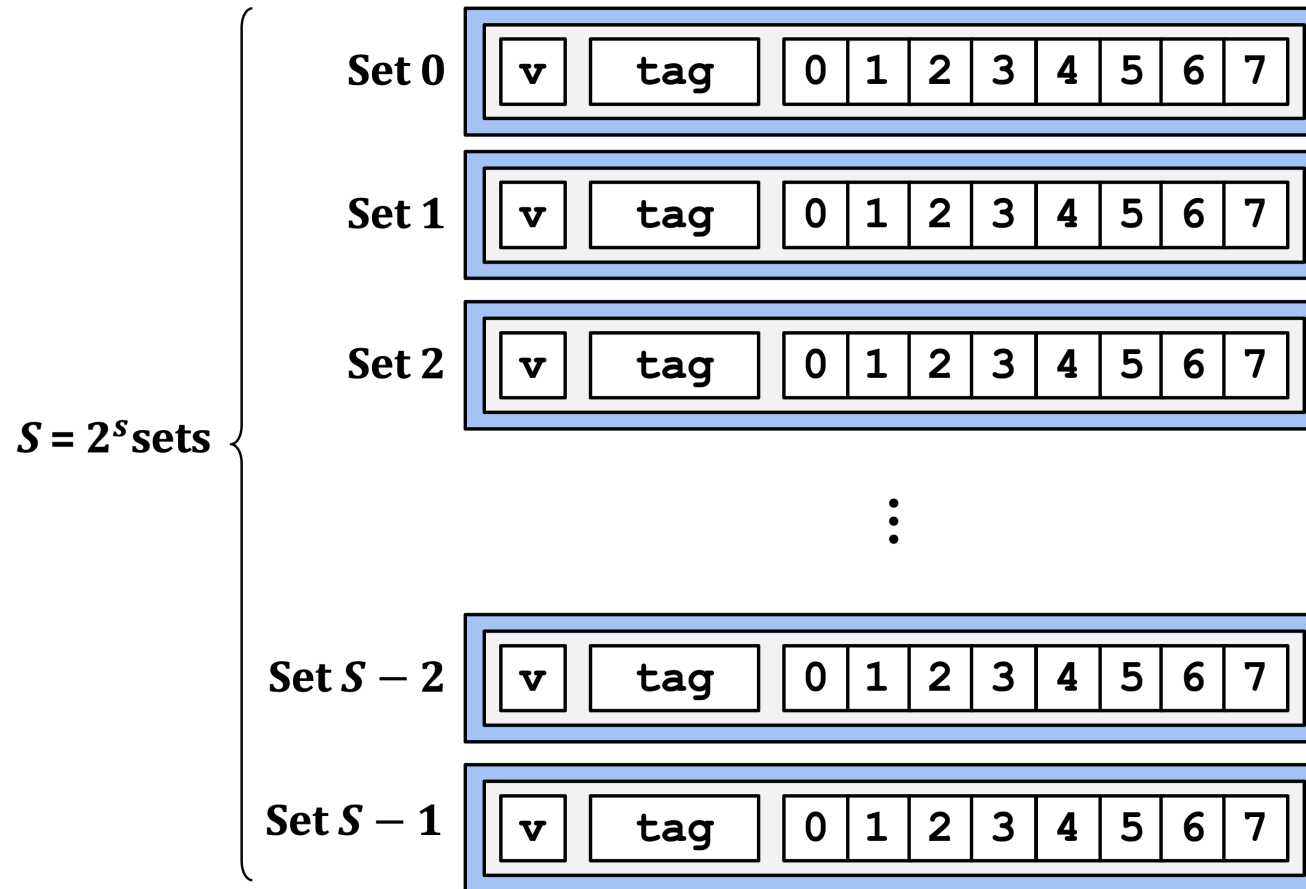
Background: Memory Address Format



- Block offset: b bits
- Set index: s bits
- Tag: $m - (s + b)$ bits

Background: Direct Mapped Cache

- Each cache set consists of a single cache line

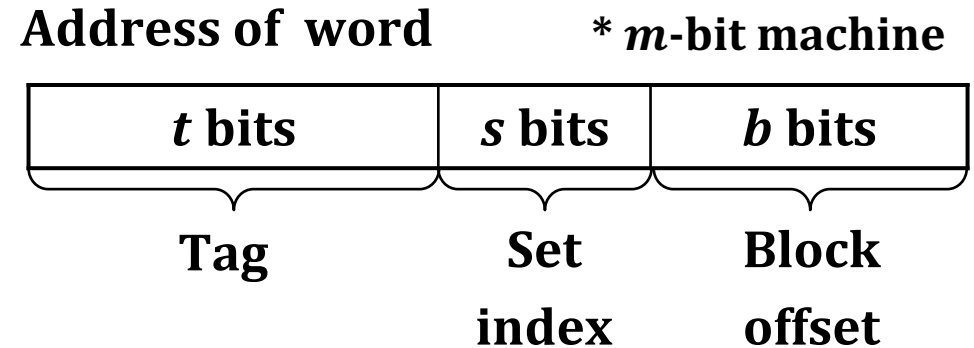
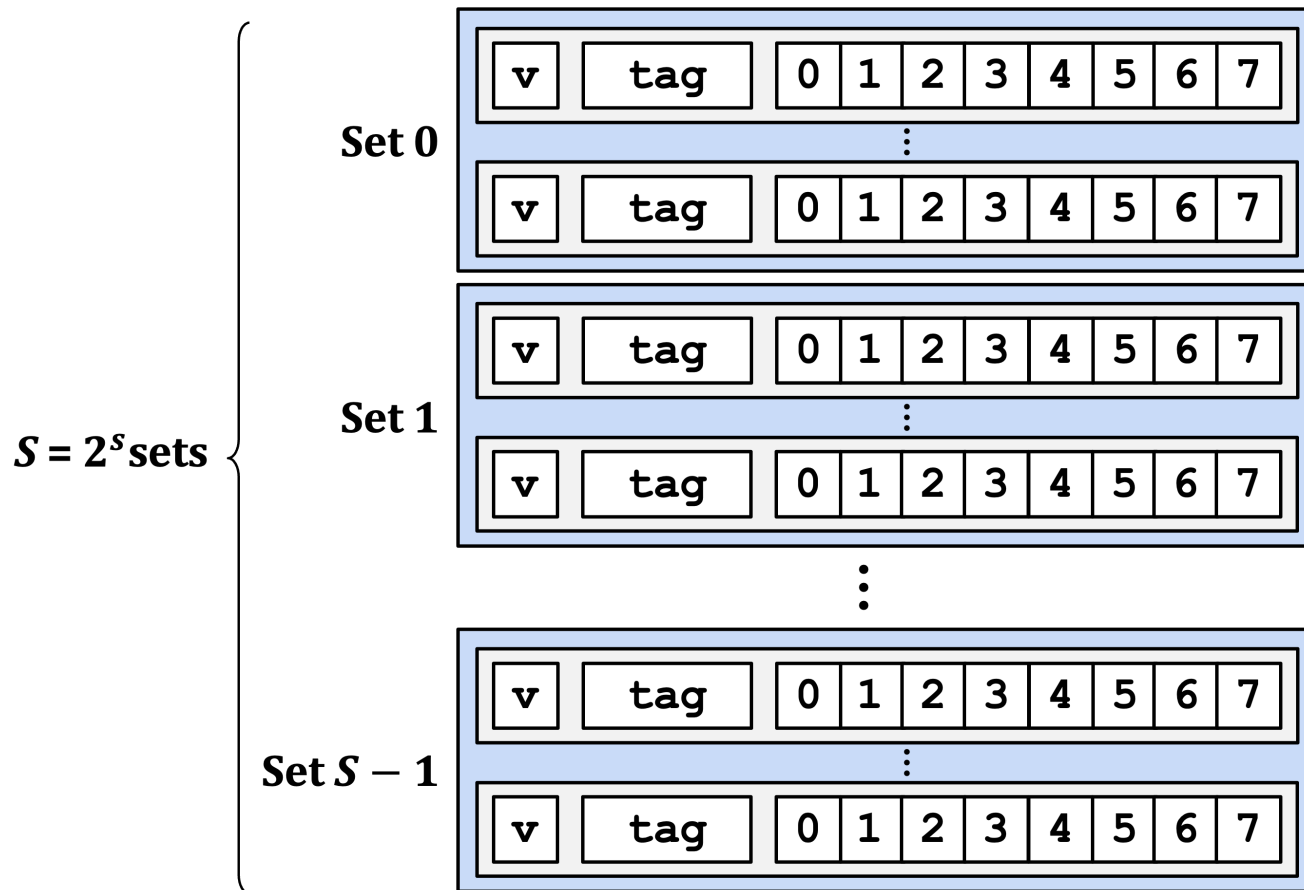


Associativity (E)	Mapping
1	Direct mapping
$1 < E < C/2^b$	Set associative mapping
$C/2^b$	Fully associative mapping

* C : Cache size

Background: Set Associative Mapped Cache

- Each cache set consists of E (associativity) cache lines

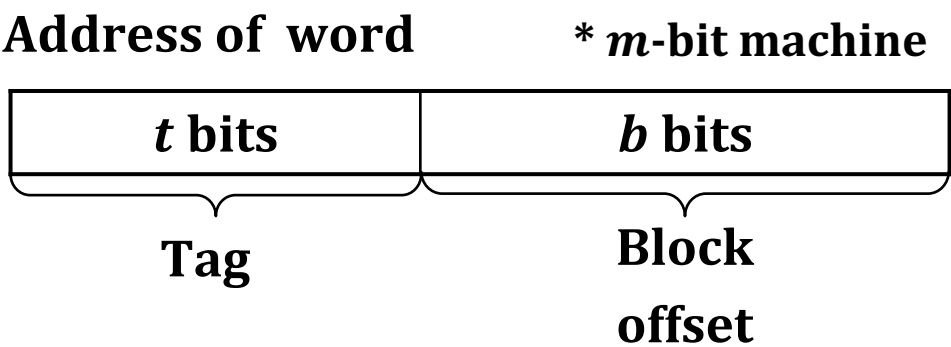
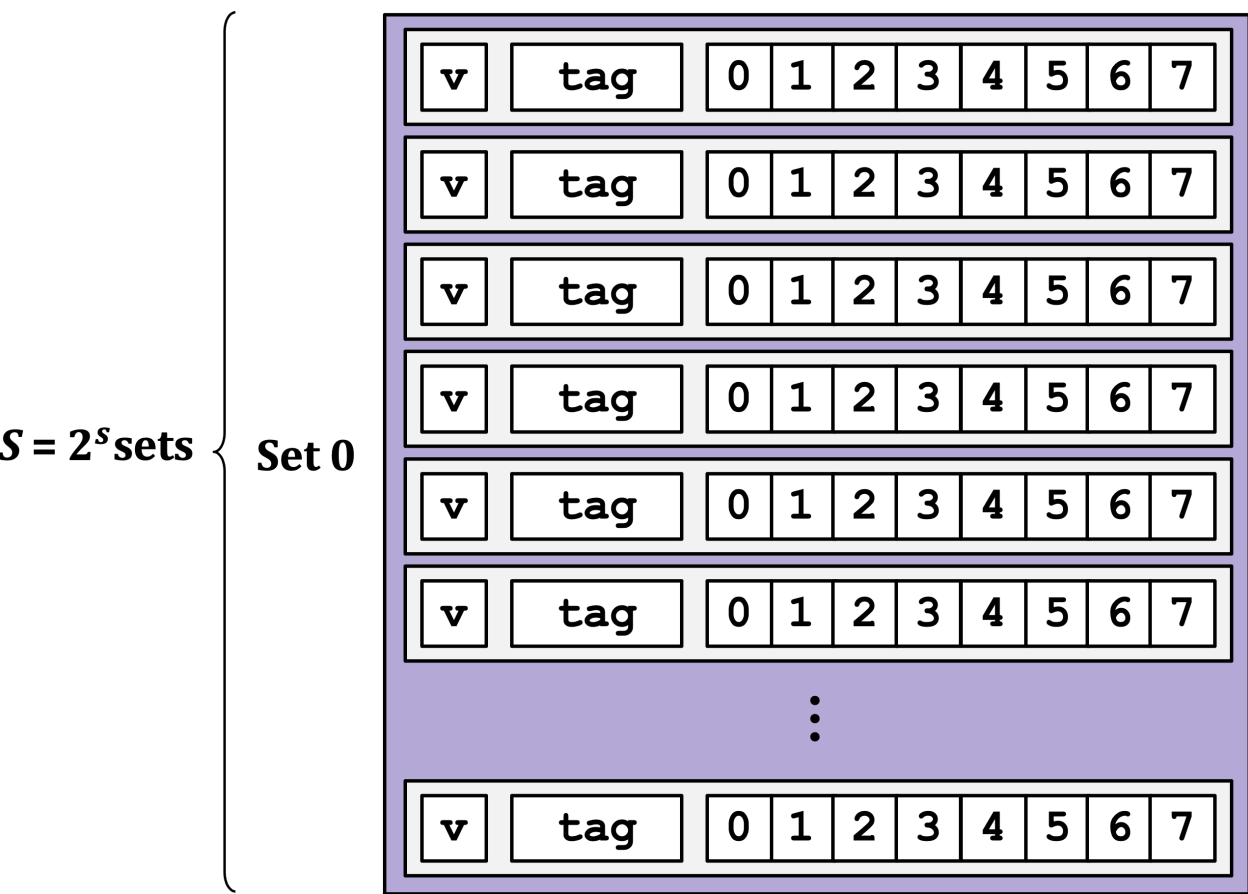


Associativity (E)	Mapping
1	Direct mapping
$1 < E < C/2^b$	Set associative mapping
$C/2^b$	Fully associative mapping

* C : Cache size

Background: Fully Associative Mapped Cache

- Cache has only one set that is composed of $C/2^b$ cache lines




No Set Index ($s = 0$ bit)

Associativity (E)	Mapping
1	Direct mapping
$1 < E < C/2^b$	Set associative mapping
$C/2^b$	Fully associative mapping

* C : Cache size


Background: Cache Hit

- Hit: An access by program to a block that is **in cache**



Background: Cache Miss

- Miss: An access by program to a block that is **not in cache**



Background: Cache Replacement Policy

- Size of cache memory is limited
 - Not all blocks from main memory can be stored in a cache
- Cache memory needs to replace **old block** with a **new block**
 - Replacement policy: Determines which block gets evicted
- Replacement policy algorithms
 - **LRU**: Replace least recently used line with the new one
 - Use counter to trace recently accessed line
 - Reference cache simulator use LRU replacement policy
 - FIFO, LIFO, etc.

Cache Lab: Overview

- Cache lab consists of Part A & B
 - Part A: Cache simulator
 - Goal: Understand the mechanism of cache memory and design cache simulator
 - Todo: Write a program simulates **cache memory access** and **counts hit/miss/eviction**
 - Part B: Efficient matrix transpose
 - Goal: Optimize matrix transpose ($A \rightarrow A^T$)
 - Todo: Write an efficient code for the highest hit ratio (i.e., minimize the cache miss)
- Submit **code files** and your **lab report (in pdf)**
 - Source code name: **[student id]_csim.c**, **[student id]_trans.c**
e.g., **20242057_csim.c**, **20242057_trans.c**
(Other formats will not be accepted)
 - Report name: **[student id].pdf** (e.g., **20242057.pdf**)

Today's Agenda

- Background
- Cache Lab Part A: Building Cache Simulator
 - Valgrind
 - File I/O APIs
 - Dynamic Allocation & Deallocation
 - Parsing Command Line Options
- Cache Lab Part B: Efficient Matrix Transpose
 - Hit Ratio
 - Matrix Multiplication

Cache Lab Part A: Building Cache Simulator

- Todo
 - Before starting Part A, read `wri teup_cachelab.pdf`
 - Write a program which simulates **cache memory access** and **count hit/miss/eviction**
- Goals
 - Understand the mechanism of cache memory and design cache simulator
- Input & Output
 - Input file is **pre-generated** by Valgrind tool
 - Output contains hit/miss/eviction for **each instruction in input file**

Valgrind

- Tool collection for debugging and profiling programs to identify memory management issues
- Usage: `$ valgrind --tool=[tool_name] [option] [program]`
 - e.g., `$ valgrind --tool=lackey ls`
- Installation (in Ubuntu)
 - `$ sudo apt install valgrind`

Tools	Function
Memcheck	Memory Profiler
Cachegrind	Cache Profiler
Callgrind	Extension of Cachegrind
Massif	Heap Profiler
Helgrind	Multi-threaded Program Debugger
Lackey	Simple Profiler

Example: Valgrind Lackey

- Lackey: Simple Valgrind tool for profiling basic operation of a program
- Other tools (e.g., massif, cachegrind) can profile various activities, such as heap usage and cache performance
 - For more details, check valgrind manual
 - <https://linux.die.net/man/1/valgrind>

```
$ valgrind --log-fd=1 --tool=lackey -v --trace-mem=yes ./your_program
```

Output logs to specified
file descriptor #

Tool name

Verbose option
(Print detailed logs)

Memory trace option
(Trace almost every **memory**
access made by the program)

File Descriptor	#
standard input	0
standard output	1
standard error	2

Valgrind Output

- Format: [operation] [address], [size]
- Operation: I (instruction load), L (data load), S (data store), M (data modify)
 - In this lab, only care about data cache activity and ignore instruction load
 - M (data modify) means one load and one store
- Address: Virtual memory address that program accessed (in HEX format)
- Size: Number of bytes used in operation

```
S 0064320c,4  
L 00603108,4  
S 00643318,4  
L 0060310c,4  
S 00643424,4  
L 00603110,4  
S 00643530,4  
L 00603114,4  
S 0064363c,4
```

File I/O APIs

- In this lab, file I/O APIs are used to read trace file (Valgrind output)
- High-level standard I/O library (e.g., `fopen()`, `fclose()`)
 - `fscanf()`: Formatted input (for reading line from trace file)
 - `fgets()`: Read a string of n bytes
- Typical usage
 - ```
fp = open("text.txt", "r");
// Do something with the file
fclose(fp);
```

# Dynamic Memory Allocation & Deallocation

- In this lab, cache memory configuration (e.g., # of cache sets, # of cache lines per set, block size) can be different by options
  - Designing simulator may need dynamic memory allocation and deallocation
- `malloc()`: Allocate consecutive dynamic memory space in heap
  - e.g., `int *p = (int*) malloc(sizeof(int) * NUM_ELEMS);`
- `free()`: De-allocates specified memory space allocated by `malloc`
  - e.g., `free(p)`

# Parsing Command Line Options

- In this lab, parsing command line options is necessary
  - e.g., help flag, verbose flag, target trace file, cache configuration
- Short options: Single-character option preceded by a single hyphen (-)
  - e.g., “ls -l”, “df -h”, “gcc -o [arg]”
- Long options: Descriptive option preceded by two hyphens (--)
  - e.g., “ls --help”, “gcc --version”, “gcc --std=c11”

# Parsing Command Line Options (Cont.)

- Option parsing functions: Included in `getopt.h` header file
  - `getopt()`: Parse short options
  - `getopt_long()`: Parse short + long options

- In this lab, only `getopt()` function is required
  - Function prototype

```
int getopt(int argc, char * const argv[], const char *optstring);
```

- Returns **ASCII value of parsed short option** (single character)
- `argc`: Number of arguments (passed through `main` function)
- `argv[]`: Array of pointers to **argument strings** (passed through `main` function)
- `optstring`: A string that specifies the options to process
  - Colon (:) means the **left option needs option argument**
  - e.g., “ho:”: “-o” option **needs** argument, but “-h” option **does not need** argument

# Parsing Command Line Options (Cont.)

- Example usage

```
int main(int argc, char *argv[]){
 int opt;
 char *outputFile;
 while((opt = getopt(argc, argv, "ho:") != -1){
 switch(opt){
 case 'h':
 print_help();
 break;
 case 'o':
 /* extern char *optarg
 * is defined in getopt.h */
 outputFile = optarg;
 print_output(outputFile);
 break;
 }
 }
 return 0;
}
```



# Programming Rules for Part A

- Your `csim.c` file **must compile without warnings** to receive credit
- You must:
  1. Design a simulator that works correctly for **arbitrary  $s$ ,  $E$ , and  $b$**
  2. **Ignore all instruction cache accesses** ("I")
  3. Call the function `printSummary()`, **at the end of your main function**
    - e.g., `printSummary(hit_count, miss_count, eviction_count);`
  4. Assume that a single memory access **never crosses block boundaries**
    - By making this assumption, you can **ignore the request sizes** in the Valgrind traces

# Today's Agenda

- Background
- Cache Lab Part A: Building Cache Simulator
  - Valgrind
  - File I/O APIs
  - Dynamic Allocation & Deallocation
  - Parsing Command Line Options
- **Cache Lab Part B: Efficient Matrix Transpose**
  - Hit Ratio
  - **Matrix Multiplication**

# Cache Lab Part B: Efficient Matrix Transpose

- Todo
  - Before starting Part B, read `wri teup_cachelab.pdf`
  - Write the efficient code for the highest hit ratio (i.e., minimize the cache miss)
- Goals
  - Optimize matrix transpose ( $A \rightarrow A^T$ )
- Cache Configuration
  - Direct mapped (E=1) cache
  - 1-kilobyte cache size
  - 32-byte (b=5) block size
  - 32 sets (s=5) in cache

# Hit Ratio

- Hit Ratio
  - Ratio of accesses that result in cache hits (# of hit / # of total access)
- Example
  - 32-byte directed mapped cache with 16-byte block size
  - Row-major order accesses

`int A[4][4]`

|             |             |             |             |
|-------------|-------------|-------------|-------------|
| <b>M</b> 0  | <b>H</b> 1  | <b>H</b> 2  | <b>H</b> 3  |
| <b>M</b> 4  | <b>H</b> 5  | <b>H</b> 6  | <b>H</b> 7  |
| <b>M</b> 8  | <b>H</b> 9  | <b>H</b> 10 | <b>H</b> 11 |
| <b>M</b> 12 | <b>H</b> 13 | <b>H</b> 14 | <b>H</b> 15 |

Cache


|           |          |          |          |
|-----------|----------|----------|----------|
| <b>8</b>  | <b>1</b> | <b>2</b> | <b>3</b> |
| <b>12</b> | <b>5</b> | <b>6</b> | <b>7</b> |

Hit Ratio = **3/4**

# Matrix Multiplication w/o Blocking

```
c = (double *) calloc(sizeof(double), n * n);

/* Multiply n x n matrices a and b */
void mmm(double *a, double *b, double *c, int n){
 int i, j, k;
 for (i = 0; i < n; i++)
 for (j = 0; j < n; j++)
 for (k = 0; k < n; k++)
 c[i * n + j] += a[i * n + k] * b[k * n + j];
}
```




# Cache Miss Analysis

- Assumptions


- Matrix elements are doubles
- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )

- First iteration

- $n/8 + n = 9n/8$  misses




*Afterwards in cache*




# Cache Miss Analysis

- Assumptions
  - Matrix elements are doubles
  - Cache block = 8 doubles
  - Cache size  $C \ll n$  (much smaller than  $n$ )

- First iteration
  - $n/8 + n = 9n/8$  misses



- Second iteration
  - Same as the first iteration




- Total misses
  - $9n/8 \times n^2 = 9/8 \times n^3$

# Matrix Multiplication w/ Blocking

```
c = (double *) calloc(sizeof(double), n * n);


/* Multiply n x n matrices a and b */
void mmm(double *a, double *b,
 double *c, int n) {
 int i, j, k;
 for (i = 0; i < n; i += B)
 for (j = 0; j < n; j += B)
 for (k = 0; k < n; k += B)
 /* B x B mini matrix multiplications */
 for (i1 = i; i1 < i + B; i1++)
 for (j1 = j; j1 < j + B; j1++)
 for (k1 = k; k1 < k + B; k1++)
 c[i1 * n + j1] +=
 a[i1 * n + k1] *
 b[k1 * n + j1];
}
```






# Cache Miss Analysis

- Assumptions

- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )
- Three blocks  fit into cache:  $3B^2 < C$


- First (block) iteration

- $B^2/8$  misses for each block
- $2n/B \times B^2/8 = nB/4$  (omitting matrix  $c$ )




# Cache Miss Analysis

- Assumptions

- Cache block = 8 doubles
- Cache size  $C \ll n$  (much smaller than  $n$ )
- Three blocks  fit into cache:  $3B^2 < C$


- First (block) iteration

- $B^2/8$  misses for each block
- $2n/B \times B^2/8 = nB/4$  (omitting matrix  $c$ )



- Second (block) iteration

- Same as the first iteration



- Total misses

- $nB/4 \times (n/B)^2 = n^3/4B$

# Blocking Summary

- No blocking:  $9/8 \times n^3$
- Blocking:  $(1/4B) \times n^3$
- Make the block size  $B$  **as large as possible**, but **do not violate**  $3B^2 < C$  condition
- Reason for dramatic difference
  - Matrix multiplication has inherent **temporal locality**
    - Input data:  $3n^2$ , computation:  $2n^3$
    - Every array elements are used  $O(n)$  times
  - **But program needs to be written properly**

# Programming Rules for Part B

- Your `trans.c` file **must compile without warnings** to receive credit
- **You are not allowed to:**
  1. **Have more than 12 local variables on the stack** at any time
    - **Define at most 12 local variables** per transpose function
    - You should also consider the case of helper functions
  2. **Side-step rule #1** by using any tricks to store more values in a single variable
    - e.g., using variables of type `long`
  3. **Define any arrays** or any variant of `malloc` (e.g., linked list)
  4. Modify array `A` (modifying array `B` is fine)
  5. Use recursion

# Cache Lab: Submission Guideline

- Due: 11/27 (Wed) 23:59 (Late submission will not be accepted)
- Submit code files and your lab report (in pdf)
  - Source code name: [student id]\_csim.c, [student id]\_trans.c  
e.g., 20242057\_csim.c, 20242057\_trans.c  
(Other formats will not be accepted)
  - Report name: [student id].pdf (e.g., 20242057.pdf)
  - A correct submission is total three files

# Cache Lab: Report Guideline

- Report
  - Attach the important parts of your code to your report
  - Explain how you built your cache simulator and optimized your matrix transpose
  - Report should not exceed 10 pages and use font Arial and font size 11pt
  - Include all references you refer to solve cache lab assignment in your report

# Cheating Policy

- You **can refer to**
  - Cache lab writeup, lab slides, and lecture slides
  - Internet sources that **do not include answers or code** related to the cache lab
    - e.g., Valgrind manual
- You **must not refer to**
  - ChatGPT with direct query **for answers or parts of a solution**
  - Code and reports from seniors who have already taken this course
  - Blogs or github repositories that **contain solution codes** (`csim.c`, `trans.c`)

# Quiz

- Go to PLMS, start the quiz
  - For fairness, **quiz will be shut down** after everyone leaves the classroom



# [CSED211] Introduction to Computer Software Systems

## Lab 6: Cache Lab

Dowon Son



**CAOS**  
COMPUTER ARCHITECTURE &  
OPERATING SYSTEMS LABORATORY

2024.11.14

# Midterm Exam Claim

- We will do claim and quiz at the same time
  - **Modifying your answer paper during the claim will be considered cheating**
- Only 10 students will do claim at the same time
  - Please check the grading criteria and think which problems to claim for others
  - From 7:45, we will prioritize the students who have the schedule after 8:00