

[CSED211] Introduction to Computer Software Systems

Lab 8: Malloc Lab

Hyeongmin Oh



CAOS
COMPUTER ARCHITECTURE &
OPERATING SYSTEMS LABORATORY

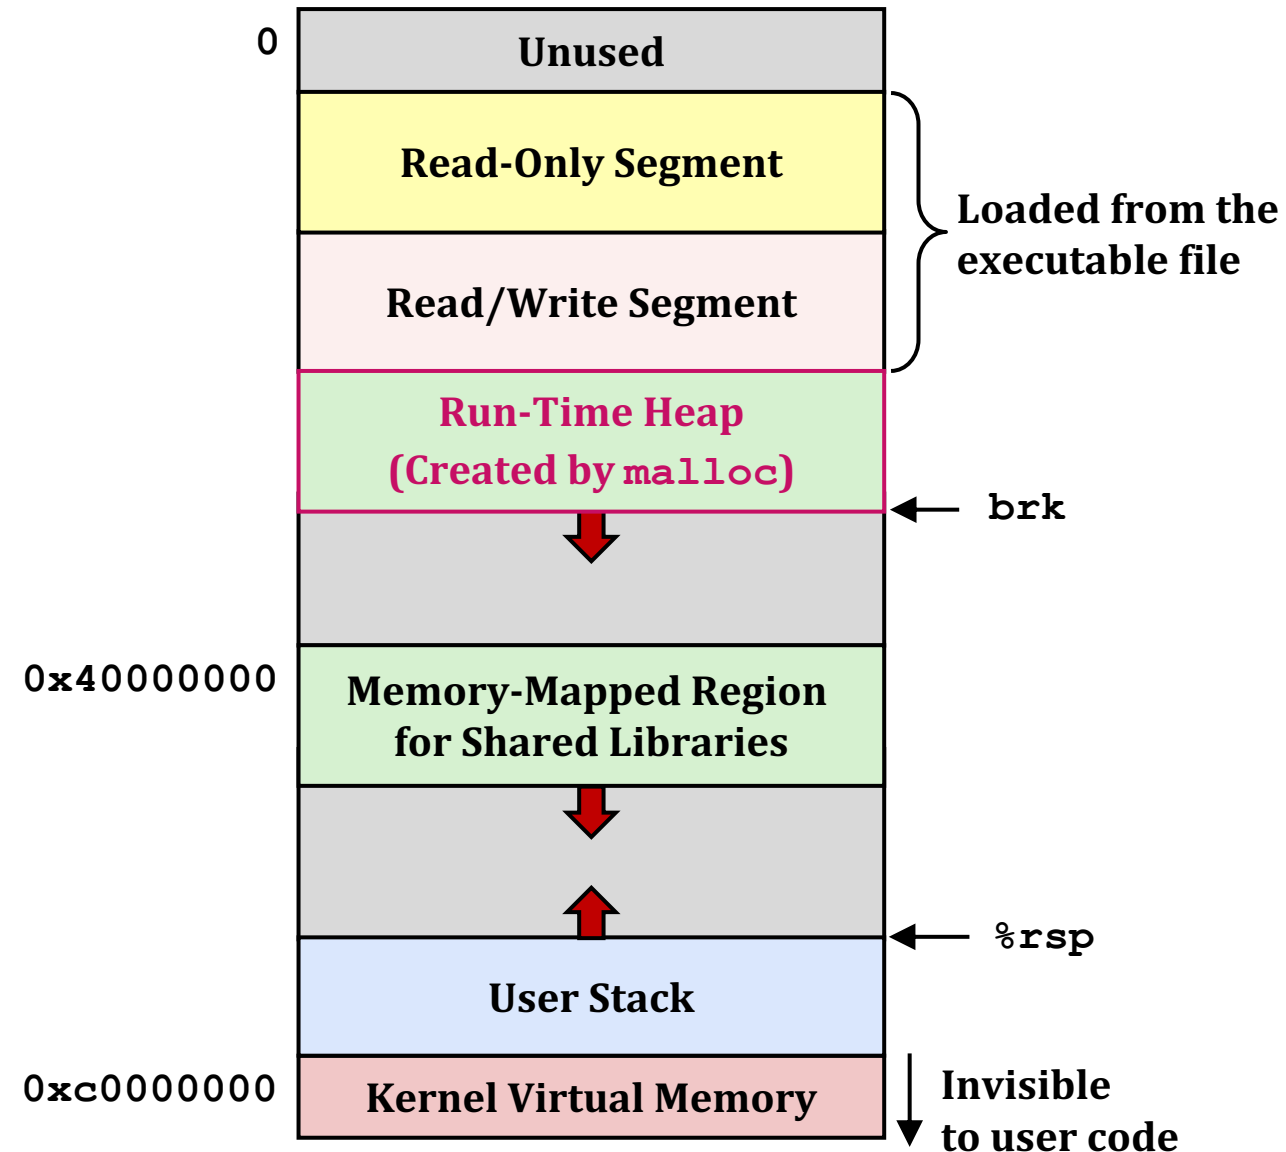
2024.12.12

Today's Agenda

- Background
 - Dynamic memory allocation
 - Example: Implicit free list
- Malloc Lab
- Quiz

Dynamic Memory Allocation

- Programmers use **dynamic memory allocators** (e.g., `malloc`) to acquire **heap memory** at runtime



Dynamic Memory Allocation (Cont.)

- Allocator maintains heap as **collection of variable sized blocks**
 - Which can be either **allocated** or **free**
- There are two types of allocator
 - **Explicit allocator**: Application allocates and frees space
 - e.g., `malloc` and `free` in C
 - **Implicit allocator**: Application only allocates, but does **not** free space
 - e.g., garbage collection in Java

Dynamic Memory Allocation (Cont.)

- Allocator maintains heap as **collection of variable sized blocks**
 - Which can be either **allocated** or **free**
- There are two types of allocator
 - **Explicit allocator**: Application allocates and frees space
 - e.g., `malloc` and `free` in C
 - **Implicit allocator**: Application only allocates, but does not free space
 - e.g., garbage collection in Java
- In this lab, we will focus on **explicit allocator**

Explicit Allocator in C (malloc package)

- Included in `stdlib.h`
- `void *malloc(size_t size)`
 - Allocate `size` bytes, and return a pointer to the address of the allocated object
- `void free(void *ptr)`
 - Free the memory space pointed by `ptr`
- `void *realloc(void *ptr, size_t size)`
 - Change the size of a previously allocated block into `size` bytes
- For more detail, check <https://linux.die.net/man/3/malloc>

Explicit Allocator in C (malloc package)

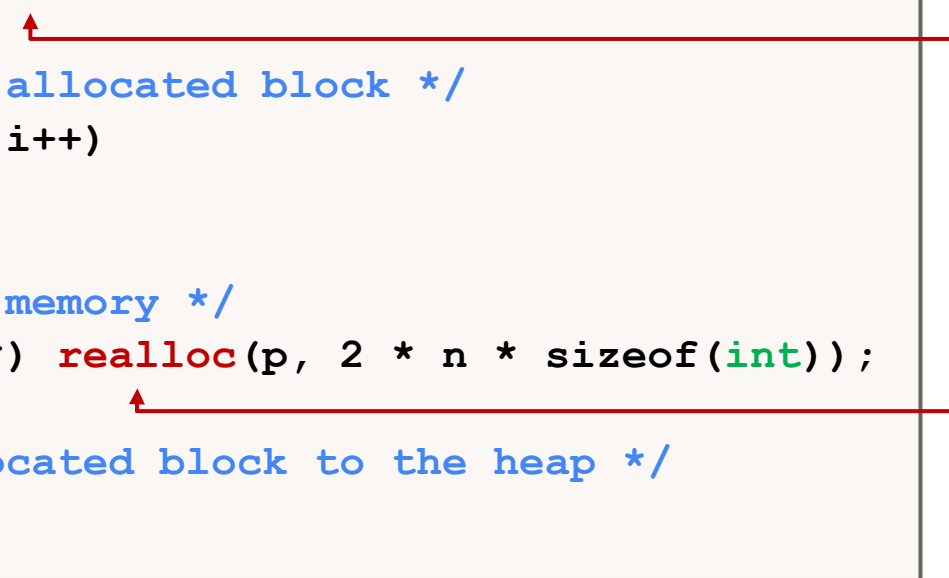
```
#include <stdio.h>
#include <stdlib.h>

void foo(int n){
    int i, *p, *p_new;
    /* Allocate a block of n ints */
    p = (int *) malloc(n * sizeof(int));

    /* Initialize allocated block */
    for(i=0; i<n; i++)
        p[i] = i;

    /* Reallocate memory */
    p_new = (int *) realloc(p, 2 * n * sizeof(int));

    /* Return allocated block to the heap */
    free(p);
}
```



We will implement
these functions!

Conditions for Good Memory Allocator

- To implement a good memory allocator, we should consider two things:
 - Throughput and memory utilization
 - They are often in a tradeoff relationship

Performance Goal: Throughput

- **Throughput**
 - The number of completed requests per unit time
 - e.g., 5,000 `malloc` calls and 5,000 `free` calls in 10 seconds
→ throughput: 1,000 operations per second
- **To maximize throughput**, the allocator should
 - Find a free memory block efficiently
 - Release an allocated memory block quickly

Performance Goal: Peak Memory Utilization


- Peak memory utilization
 - For a given sequence of malloc and free requests $R_0, R_1, \dots, R_k, \dots, R_{N-1}$
 - Aggregate payload P_k : Sum of currently allocated payloads after completing R_k
 - Current heap size H_k : Monotonically increase when allocator uses `sbrk()`
 - Peak memory utilization U_k after completing $(k + 1)$ requests: $U_k = \max_{i \leq k}(P_i)/H_k$
- To maximize peak memory utilization, allocator should avoid fragmentation

Fragmentation

- There are two kinds of fragmentation
 - Internal fragmentation
 - External fragmentation

Internal Fragmentation

- Occurs if the payload is smaller than the block size




- Caused by
 - Overhead of maintaining heap data structures (e.g., boundary tag, pointer)
 - Padding for alignment purposes
 - Explicit policy decisions (e.g., return a big block to satisfy a small request)

External Fragmentation


- Occurs when allocating blocks and freeing the middle one, leaving small pieces of free space that may be too small to reuse

`#define SZ sizeof(long)`




Free Block Management Schemes

- Method 1: **Implicit list** using length – links all blocks
 - Needs to tag each block as allocated or free




- Method 2: **Explicit list** among the free blocks using pointers
 - Needs space for pointers



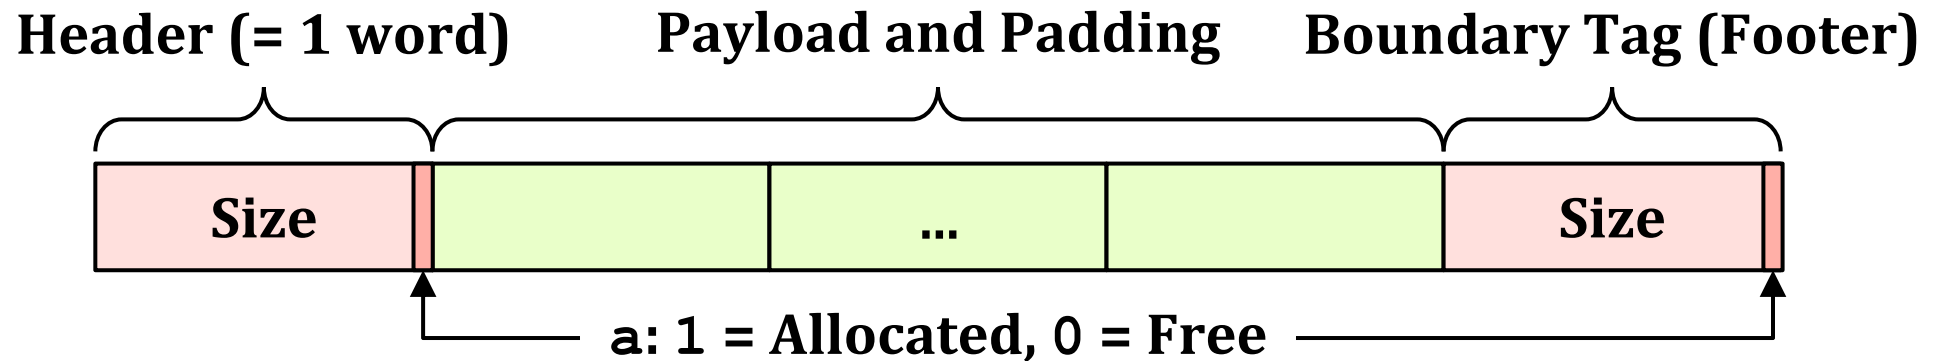
- Method 3: **Segregated free list** - different free lists for different size classes
- Method 4: **Blocks sorted by size** - can use a balanced tree (e.g., red-black tree) with pointers within each free block, and the length used as a key

Implicit Free List

- **Implicit: No explicit pointer** to reference adjacent blocks



- Instead, **location of the adjacent block can be inferred** using boundary tags
 - Two boundary tags (header, footer) are needed to traverse bidirectional



Implicit Free List: Design Space

- There are multiple decision points when implementing implicit free list
- Placement policy
 - During `malloc`, what free block should be used for allocation
- Splitting policy
 - After `malloc`, split the remaining free space as a separate block or not
- Coalescing policy
 - After `free`, merge adjacent free block immediately or not


Implicit Free List: Design Space

- Placement policy
 - First fit
 - Searches the list from the beginning and chooses the **first** free block that fits
 - Next fit
 - Like first fit, but searches the list **starting where previous search finished**
 - Best fit
 - Searches the list and chooses the **best** free block that fits with the fewest bytes left

Implicit Free List: Design Space

- **Splitting policy**

- Since allocated space might be smaller than the free space, it might need to split the free block



- **Coalescing policy**

- **Immediate coalescing**: Coalesce each time `free` is called
- **Deferred coalescing**: Try to improve performance of `free` by deferring coalescing until needed

Implicit Free List: Implementation

- You can implement implicit free list by referring CS:APP textbook
 - Placement policy: First-fit
 - Splitting policy: Enable splitting
 - Coalescing policy: Immediate coalescing
 - Check pp.883-897 for more detail

Results for mm malloc:

trace	valid	util	ops	secs	Kops
0	yes	99%	5694	0.022506	253
1	yes	100%	5848	0.019721	297
...					
24	yes	89%	12	0.000000	40000
25	yes	89%	12	0.000000	40000
Total		75%	209279	2.882481	73

Perf index = 45 (util) + 5 (thru) = 50/100

Implicit Free List: Performance Analysis

- Moderate memory utilization (45/60)

Perf index = 45 (util) + 5 (thru) = 50/100

- Internal fragmentation: Good

- Small data structure overhead compared to other free block management schemes

- External fragmentation: Can be bad

- First-fit does not care about remaining free block capacity
- Fragmented free block keeps accumulated

- Low throughput (5/40)

- If tiny `malloc()` requests are dominant, small blocks are accumulated at the front

- Since implicit free list also traverses allocated block when finding available free block, throughput can be significantly decreased

- Even though `coalescing()` is constant-time procedure, immediate coalescing can degrade throughput

- If free block is used for the subsequent request, coalescing is unnecessary

Today's Agenda

- Background
- **Malloc Lab**
- Quiz

Malloc Lab: Assignment

- Goal: Gain hands-on experience about dynamic memory allocation
- In this lab, you will implement five functions declared in `mm.c`
 - `mm_init()`
 - `mm_malloc()`
 - `mm_free()`
 - `mm_realloc()`
 - `mm_check()` (for debugging, disable when you submit)
- Complete `mm.c` and run:
 - `$ make && ./mdriver -V`

Malloc Lab: Submission Guideline

- Due: 12/23 (Mon) 23:59
 - Late submission will not be accepted
- Two files must be uploaded to PLMS: mm.c and lab report (in pdf)
- Submission format (0 point if violated)
 - Code name: [student id]_mm.c (e.g., 2023xxxx_mm.c)
 - Report name: [student id].pdf (e.g., 2023xxxx.pdf)
- No page, font limit on lab report
 - Show every work you did!

Malloc Lab: Evaluation

- **Score evaluation:** Quiz (10%) + Code (40%) + Report (50%)
 - Code evaluation criteria
 - Performance index will be translated based on the score table

Perf. Index	<50	50~59	60~69	70~79	80~89	90~100
Score	10 pts	20 pts	25 pts	30 pts	35 pts	40 pts

- Report evaluation criteria

function	<code>mm_malloc()</code>	<code>mm_free()</code>	<code>mm_realloc()</code>	<code>mm_init()</code>	<code>mm_check()</code>
Score	15 pts	15 pts	10 pts	5 pts	5 pts

Malloc Lab: Coding Rules (0 point if violated)

- Your code should compile & run successfully
- Return pointers of your allocator must be **8-byte aligned**
 - Malloc lab assumes 32-bit system, so double word is 8-byte
- **Do not**
 - Change the interface (function header) in `mm.c`
 - Use any memory-management related library or system calls
 - e.g., `malloc`, `free`, `realloc`, `calloc`, `sbrk`, `brk`
 - Do not implement `mm_malloc` using `malloc`, `mm_free` using `free`, ...
 - `memcpy` is okay
 - Define any global/static compound data structure in `mm.c`
 - e.g., arrays, structs, trees, lists, ...

Malloc Lab: Lab Report Guideline

- Please clarify the type of memory allocator you implemented
 - e.g., implicit free list, explicit free list, segregated free list, ...
- For each function, please provide
 1. Screenshot of your code
 2. How does your code work (detailed comment is enough)
- Provide the rationale behind your throughput / memory utilization score
 - e.g., why your memory allocator shows poor memory utilization, why your memory allocator shows high throughput, ...

Cheating Policy

- You **can refer**
 - Textbook source code
 - Malloc lab writeup, lab slides, lecture slides
 - Internet sources that **doesn't involve direct answers or codes** about malloc lab
- You **should not refer**
 - ChatGPT
 - Code and report from a senior who has already taken this course
 - Blogs or github repository that contain solution codes
 - Every other references that violate POSTECHIAN's Honor

Quiz

- Go to the PLMS, and start the Quiz!